

# Moving from Latch

Everything you run on Latch runs on Carolina Cloud at a fraction of the cost — and we'll migrate you, free.

## SUMMARY

Your pipelines are Nextflow and your files are ordinary S3 storage, so there is no lock-in to unwind — moving is mostly copying data. Here's how it goes:

- 1 Email us**  
 Send your rough data size and the pipelines you run to **derek@carolinacloud.io**
- 2 We move your data**  
 We copy it into your new bucket for you — big transfers run unattended overnight. Buckets you already own keep working as they are.
- 3 We set up your workspace**  
 Instances, teammates and roles, and your first pipeline running end-to-end. You watch it work before committing to anything.
- 4 You switch when it's proven**  
 You pay nothing until you're fully moved — up to **\$5,000 in migration credits** covers the switch. No migration is too small.

Carolina Cloud is owner-operated and profitable, and we run our own hardware. No egress fees, ever — downloading your data, sharing across clouds, or leaving later is free.

## CONCEPT MAP

| On Latch                                                 | On Carolina Cloud                                                                                                                                                                     |
|----------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Latch Data</b> (latch:// paths, file browser)         | <b>Buckets</b> — S3-compatible storage with a built-in file browser, usable from any S3 client                                                                                        |
| <b>Mounted data</b> (Latch Data on Pods, S3/GCP buckets) | Your bucket is mounted at ~/personalbucket on every instance (plus ~/orgbucket for your org's shared bucket); S3/GCP buckets you already own keep working, or we copy them in for you |
| <b>Workflows</b> (Nextflow via the Latch SDK)            | <b>Native Nextflow</b> — run plain nf-core or your own pipelines with no SDK, no registration step, no wrapper code, at roughly 10–20x lower compute cost                             |
| <b>Pods</b> (JupyterLab, RStudio)                        | <b>Instances</b> — JupyterLab, RStudio Server, Marimo notebooks, and SSH containers; live-resizable without a restart, and clonable for a teammate                                    |
| <b>Pod GPUs</b> (A10G, T4, V100, L40s)                   | Blackwell GPUs, just added to the fleet — available in notebooks and compute containers                                                                                               |
| <b>Workspaces &amp; Organizations</b>                    | <b>Organizations</b> — shared credit pool, admin/member roles, per-member spend caps, required 2FA                                                                                    |
| <b>Registry</b> (sample sheets, metadata)                | Samplesheets live as CSVs in your bucket — which is what Nextflow pipelines consume anyway                                                                                            |
| <b>Plots</b> (interactive dashboards)                    | <b>Marimo</b> — reactive notebooks that deploy as interactive apps                                                                                                                    |

## PRICING VS LATCH

|                                                          |                                                         |                                                        |                                                              |
|----------------------------------------------------------|---------------------------------------------------------|--------------------------------------------------------|--------------------------------------------------------------|
| <b>20×</b><br>cheaper CPU<br>\$0.21 → \$0.01 per vCPU-hr | <b>2×</b><br>cheaper RAM<br>\$0.01 → \$0.005 per GiB-hr | <b>2.5×</b><br>cheaper storage<br>\$23 → \$9 per TB-mo | <b>\$0</b><br>egress fees<br>downloads & sharing always free |
|----------------------------------------------------------|---------------------------------------------------------|--------------------------------------------------------|--------------------------------------------------------------|

Both platforms bill per second, 1 credit = \$1 · rates from latch.bio/pricing, Sep 2026 · console.carolinacloud.io